

# CMTH642 Assignment 3

Geoffrey Clark

March 27, 2018

## Data Prep

Preparation: The dataset is related to white Portuguese “Vinh o Verde” wine. For more info: <https://archive.ics.uci.edu/ml/datasets/Wine+Quality>

Import to R the following file: <http://archive.ics.uci.edu/ml/machine-learning-databases/wine-quality/winequality-white.csv>

```
# Download remote content
# wine.df <- read.csv("http://archive.ics.uci.edu/ml/machine-learning-databases/wine-quality/winequality-white.csv")
wine.df <- read.csv("winequality-white.csv", header=T, sep=";")
```

## QUESTIONS

1. Check data characteristics. Is there missing data?

```
sum(is.na(wine.df)) #0
```

```
## [1] 0
```

```
str(wine.df)
```

```
## 'data.frame': 4898 obs. of 12 variables:
## $ fixed.acidity : num 7 6.3 8.1 7.2 7.2 8.1 6.2 7 6.3 8.1 ...
## $ volatile.acidity : num 0.27 0.3 0.28 0.23 0.23 0.28 0.32 0.27 0.3 0.22 ...
## $ citric.acid : num 0.36 0.34 0.4 0.32 0.32 0.4 0.16 0.36 0.34 0.43 ...
## $ residual.sugar : num 20.7 1.6 6.9 8.5 8.5 6.9 7 20.7 1.6 1.5 ...
## $ chlorides : num 0.045 0.049 0.05 0.058 0.058 0.05 0.045 0.045 0.049 0.044 ...
## $ free.sulfur.dioxide : num 45 14 30 47 47 30 30 45 14 28 ...
## $ total.sulfur.dioxide: num 170 132 97 186 186 97 136 170 132 129 ...
## $ density : num 1.001 0.994 0.995 0.996 0.996 ...
## $ pH : num 3 3.3 3.26 3.19 3.19 3.26 3.18 3 3.3 3.22 ...
## $ sulphates : num 0.45 0.49 0.44 0.4 0.4 0.44 0.47 0.45 0.49 0.45 ...
## $ alcohol : num 8.8 9.5 10.1 9.9 9.9 10.1 9.6 8.8 9.5 11 ...
## $ quality : int 6 6 6 6 6 6 6 6 6 6 ...
```

```
sapply(wine.df, function(x) sum(is.na(x))) # 0 NAs
```

```
##      fixed.acidity    volatile.acidity    citric.acid
##           0           0           0
##      residual.sugar    chlorides    free.sulfur.dioxide
##           0           0           0
## total.sulfur.dioxide    density    pH
##           0           0           0
##      sulphates    alcohol    quality
##           0           0           0
```

```
sapply(wine.df, function(x) sum(is.null(x))) # 0 NULLs
```

```
##      fixed.acidity    volatile.acidity    citric.acid
##           0           0           0
##      residual.sugar    chlorides    free.sulfur.dioxide
```

```
##          0          0          0
## total.sulfur.dioxide      density      pH
##          0          0          0
##          sulphates      alcohol      quality
##          0          0          0
```

2. What is the correlation between the attributes other than wine quality?

```
# For Visualizing Correlation:
# install.packages("corrplot")
library(corrplot)
```


```
## corrplot 0.84 loaded
```

```
(wine.cor <- cor(wine.df[, -12]))
```

```
##          fixed.acidity volatile.acidity citric.acid
## fixed.acidity      1.00000000      -0.02269729  0.28918070
## volatile.acidity    -0.02269729      1.00000000 -0.14947181
## citric.acid         0.28918070     -0.14947181  1.00000000
## residual.sugar      0.08902070      0.06428606  0.09421162
## chlorides           0.02308564      0.07051157  0.11436445
## free.sulfur.dioxide -0.04939586     -0.09701194  0.09407722
## total.sulfur.dioxide 0.09106976      0.08926050  0.12113080
## density             0.26533101      0.02711385  0.14950257
## pH                 -0.42585829     -0.03191537 -0.16374821
## sulphates          -0.01714299     -0.03572815  0.06233094
## alcohol            -0.12088112      0.06771794 -0.07572873
##          residual.sugar  chlorides free.sulfur.dioxide
## fixed.acidity      0.08902070  0.02308564      -0.0493958591
## volatile.acidity    0.06428606  0.07051157      -0.0970119393
## citric.acid         0.09421162  0.11436445      0.0940772210
## residual.sugar      1.00000000  0.08868454      0.2990983537
## chlorides           0.08868454  1.00000000      0.1013923521
## free.sulfur.dioxide 0.29909835  0.10139235      1.0000000000
## total.sulfur.dioxide 0.40143931  0.19891030      0.6155009650
## density             0.83896645  0.25721132      0.2942104109
## pH                 -0.19413345 -0.09043946     -0.0006177961
## sulphates          -0.02666437  0.01676288      0.0592172458
## alcohol            -0.45063122 -0.36018871     -0.2501039415
##          total.sulfur.dioxide      density      pH
## fixed.acidity      0.091069756  0.26533101 -0.4258582910
## volatile.acidity    0.089260504  0.02711385 -0.0319153683
## citric.acid         0.121130798  0.14950257 -0.1637482114
## residual.sugar      0.401439311  0.83896645 -0.1941334540
## chlorides           0.198910300  0.25721132 -0.0904394560
## free.sulfur.dioxide 0.615500965  0.29421041 -0.0006177961
## total.sulfur.dioxide 1.000000000  0.52988132  0.0023209718
## density             0.529881324  1.00000000 -0.0935914935
## pH                 0.002320972 -0.09359149  1.0000000000
## sulphates           0.134562367  0.07449315  0.1559514973
## alcohol            -0.448892102 -0.78013762  0.1214320987
##          sulphates      alcohol
## fixed.acidity      -0.01714299 -0.12088112
## volatile.acidity    -0.03572815  0.06771794
## citric.acid         0.06233094 -0.07572873
```

```
## residual.sugar      -0.02666437 -0.45063122
## chlorides           0.01676288 -0.36018871
## free.sulfur.dioxide 0.05921725 -0.25010394
## total.sulfur.dioxide 0.13456237 -0.44889210
## density             0.07449315 -0.78013762
## pH                 0.15595150  0.12143210
## sulphates           1.00000000 -0.01743277
## alcohol             -0.01743277  1.00000000
```

```
corrplot(wine.cor, method="number") # I did this because I appreciate the visual
```



```
# highest correlation:
```

```
# density & residual.sugar: 0.84
```

```
# density & alcohol: -0.78
```

```
# I found a cool SO thread on how to do this programmatically:
```

```
# https://stackoverflow.com/questions/7074246/show-correlations-as-an-ordered-list-not-as-a-large-matrix
```

```
# as.data.frame(as.table(wine.cor))
```

```
# install.packages("reshape")
```

```
library(reshape) # includes melt()
```

```
wine.cor.list <- wine.cor
```

```
wine.cor.list[wine.cor.list == 1] <- NA
```

```
wine.cor.list <- na.omit(melt(wine.cor.list))
```

```
wine.cor.list[order(-abs(wine.cor.list$value)),]
```

##	X1	X2	value
## 41	density	residual.sugar	0.8389664549
## 81	residual.sugar	density	0.8389664549
## 88	alcohol	density	-0.7801376214
## 118	density	alcohol	-0.7801376214
## 62	total.sulfur.dioxide	free.sulfur.dioxide	0.6155009650
## 72	free.sulfur.dioxide	total.sulfur.dioxide	0.6155009650
## 74	density	total.sulfur.dioxide	0.5298813239
## 84	total.sulfur.dioxide	density	0.5298813239
## 44	alcohol	residual.sugar	-0.4506312220
## 114	residual.sugar	alcohol	-0.4506312220
## 77	alcohol	total.sulfur.dioxide	-0.4488921021
## 117	total.sulfur.dioxide	alcohol	-0.4488921021
## 9	pH	fixed.acidity	-0.4258582910
## 89	fixed.acidity	pH	-0.4258582910
## 40	total.sulfur.dioxide	residual.sugar	0.4014393112
## 70	residual.sugar	total.sulfur.dioxide	0.4014393112
## 55	alcohol	chlorides	-0.3601887121
## 115	chlorides	alcohol	-0.3601887121
## 39	free.sulfur.dioxide	residual.sugar	0.2990983537
## 59	residual.sugar	free.sulfur.dioxide	0.2990983537
## 63	density	free.sulfur.dioxide	0.2942104109
## 83	free.sulfur.dioxide	density	0.2942104109
## 3	citric.acid	fixed.acidity	0.2891806977
## 23	fixed.acidity	citric.acid	0.2891806977
## 8	density	fixed.acidity	0.2653310138
## 78	fixed.acidity	density	0.2653310138
## 52	density	chlorides	0.2572113204
## 82	chlorides	density	0.2572113204
## 66	alcohol	free.sulfur.dioxide	-0.2501039415
## 116	free.sulfur.dioxide	alcohol	-0.2501039415
## 51	total.sulfur.dioxide	chlorides	0.1989102996
## 71	chlorides	total.sulfur.dioxide	0.1989102996
## 42	pH	residual.sugar	-0.1941334540
## 92	residual.sugar	pH	-0.1941334540
## 31	pH	citric.acid	-0.1637482114
## 91	citric.acid	pH	-0.1637482114
## 98	sulphates	pH	0.1559514973
## 108	pH	sulphates	0.1559514973
## 30	density	citric.acid	0.1495025706
## 80	citric.acid	density	0.1495025706
## 14	citric.acid	volatile.acidity	-0.1494718106
## 24	volatile.acidity	citric.acid	-0.1494718106
## 76	sulphates	total.sulfur.dioxide	0.1345623669
## 106	total.sulfur.dioxide	sulphates	0.1345623669
## 99	alcohol	pH	0.1214320987
## 119	pH	alcohol	0.1214320987
## 29	total.sulfur.dioxide	citric.acid	0.1211307977
## 69	citric.acid	total.sulfur.dioxide	0.1211307977
## 11	alcohol	fixed.acidity	-0.1208811232
## 111	fixed.acidity	alcohol	-0.1208811232
## 27	chlorides	citric.acid	0.1143644484
## 47	citric.acid	chlorides	0.1143644484
## 50	free.sulfur.dioxide	chlorides	0.1013923521

## 60	chlorides	free.sulfur.dioxide	0.1013923521
## 17	free.sulfur.dioxide	volatile.acidity	-0.0970119393
## 57	volatile.acidity	free.sulfur.dioxide	-0.0970119393
## 26	residual.sugar	citric.acid	0.0942116243
## 36	citric.acid	residual.sugar	0.0942116243
## 28	free.sulfur.dioxide	citric.acid	0.0940772210
## 58	citric.acid	free.sulfur.dioxide	0.0940772210
## 86	pH	density	-0.0935914935
## 96	density	pH	-0.0935914935
## 7	total.sulfur.dioxide	fixed.acidity	0.0910697562
## 67	fixed.acidity	total.sulfur.dioxide	0.0910697562
## 53	pH	chlorides	-0.0904394560
## 93	chlorides	pH	-0.0904394560
## 18	total.sulfur.dioxide	volatile.acidity	0.0892605036
## 68	volatile.acidity	total.sulfur.dioxide	0.0892605036
## 4	residual.sugar	fixed.acidity	0.0890207014
## 34	fixed.acidity	residual.sugar	0.0890207014
## 38	chlorides	residual.sugar	0.0886845359
## 48	residual.sugar	chlorides	0.0886845359
## 33	alcohol	citric.acid	-0.0757287301
## 113	citric.acid	alcohol	-0.0757287301
## 87	sulphates	density	0.0744931485
## 107	density	sulphates	0.0744931485
## 16	chlorides	volatile.acidity	0.0705115715
## 46	volatile.acidity	chlorides	0.0705115715
## 22	alcohol	volatile.acidity	0.0677179428
## 112	volatile.acidity	alcohol	0.0677179428
## 15	residual.sugar	volatile.acidity	0.0642860601
## 35	volatile.acidity	residual.sugar	0.0642860601
## 32	sulphates	citric.acid	0.0623309403
## 102	citric.acid	sulphates	0.0623309403
## 65	sulphates	free.sulfur.dioxide	0.0592172458
## 105	free.sulfur.dioxide	sulphates	0.0592172458
## 6	free.sulfur.dioxide	fixed.acidity	-0.0493958591
## 56	fixed.acidity	free.sulfur.dioxide	-0.0493958591
## 21	sulphates	volatile.acidity	-0.0357281469
## 101	volatile.acidity	sulphates	-0.0357281469
## 20	pH	volatile.acidity	-0.0319153683
## 90	volatile.acidity	pH	-0.0319153683
## 19	density	volatile.acidity	0.0271138455
## 79	volatile.acidity	density	0.0271138455
## 43	sulphates	residual.sugar	-0.0266643659
## 103	residual.sugar	sulphates	-0.0266643659
## 5	chlorides	fixed.acidity	0.0230856437
## 45	fixed.acidity	chlorides	0.0230856437
## 2	volatile.acidity	fixed.acidity	-0.0226972901
## 12	fixed.acidity	volatile.acidity	-0.0226972901
## 110	alcohol	sulphates	-0.0174327719
## 120	sulphates	alcohol	-0.0174327719
## 10	sulphates	fixed.acidity	-0.0171429850
## 100	fixed.acidity	sulphates	-0.0171429850
## 54	sulphates	chlorides	0.0167628837
## 104	chlorides	sulphates	0.0167628837
## 75	pH	total.sulfur.dioxide	0.0023209718

```
## 95 total.sulfur.dioxide      pH  0.0023209718
## 64                pH  free.sulfur.dioxide -0.0006177961
## 94  free.sulfur.dioxide      pH -0.0006177961
```


3. Graph the frequency distribution of wine quality .

```
hist(wine.df$quality, freq=T, breaks=seq(2.5,9.5,1))
```



4. Reduce the levels of rating for quality to three levels as high, medium and low.

```
wine.df$quality <- cut(wine.df$quality, 3, labels=c("Low","Medium","High"))
plot(wine.df$quality) # unequal distribution of class variables (potential class imbalance problem)
```



5. Normalize the data set.

```
# scale() returns a matrix. Also requires numerical values (exclude class attribute (factor): column 12
wine.df.norm <- as.data.frame(scale(wine.df[, -12]))
wine.df.norm <- cbind(wine.df.norm, wine.df[, 12]) # Re-add class attribute
names(wine.df.norm)[12] <- names(wine.df)[12] # Re-name the class attribute
```

6. Divide the data to training and testing groups.

```
wine.train_index <- sample(1:nrow(wine.df.norm), 0.7*nrow(wine.df.norm))
wine.train <- wine.df.norm[wine.train_index,]
wine.test <- wine.df.norm[-wine.train_index,]
(nrow(wine.df) == (nrow(wine.train) + nrow(wine.test))) # Did we include all observations?

## [1] TRUE
```

7. Use the KNN algorithm to predict the quality of wine using its attributes.

```
library(class) # needed for knn()

##
## Attaching package: 'class'
## The following object is masked from 'package:reshape':
##
##      condense

# I chose to reference the test group directly in the following calls to knn().
# wine.train_labels <- wine.train$quality
# wine.test_labels <- wine.test$quality
```

```
#
# note the subsetting & removing the attribute quality (wine.test[, -12]): this
# was done because knn() doesn't accept numerical values and seems to react
# poorly to having the class attribute included in either train or test sets.

wine.knn <- list(k3=factor(), k5=factor(), k7=factor())
wine.knn$k3 <- knn(train=wine.train[, -12], test=wine.test[, -12], cl=wine.train$quality, k=3, prob=T)
wine.knn$k5 <- knn(train=wine.train[, -12], test=wine.test[, -12], cl=wine.train$quality, k=5)
wine.knn$k7 <- knn(train=wine.train[, -12], test=wine.test[, -12], cl=wine.train$quality, k=7)
```

8. Evaluate the model performance .

```
library(gmodels) # Necessary for CrossTable()

# I prefer actual results along the top of the table: I find it easier to read
# This seems pretty standard, and is likewise on Wikipedia: https://en.wikipedia.org/wiki/Confusion\_matrix
CrossTable(x=wine.knn$k3, y=wine.test$quality, prop.chisq = F)
```

```
##
##
##      Cell Contents
## |-----|
## |                      N |
## |      N / Row Total |
## |      N / Col Total |
## |      N / Table Total |
## |-----|
##
##
## Total Observations in Table:  1470
##
##
##      | wine.test$quality
## wine.knn$k3 |      Low |      Medium |      High | Row Total |
## -----|-----|-----|-----|-----|
##      Low |      278 |      149 |           1 |      428 |
##      |      0.650 |      0.348 |      0.002 |      0.291 |
##      |      0.566 |      0.160 |      0.020 |           |
##      |      0.189 |      0.101 |      0.001 |           |
## -----|-----|-----|-----|
##      Medium |      207 |      759 |           38 |     1004 |
##      |      0.206 |      0.756 |      0.038 |      0.683 |
##      |      0.422 |      0.816 |      0.776 |           |
##      |      0.141 |      0.516 |      0.026 |           |
## -----|-----|-----|-----|
##      High |           6 |           22 |           10 |           38 |
##      |      0.158 |      0.579 |      0.263 |      0.026 |
##      |      0.012 |      0.024 |      0.204 |           |
##      |      0.004 |      0.015 |      0.007 |           |
## -----|-----|-----|-----|
## Column Total |      491 |      930 |           49 |     1470 |
##      |      0.334 |      0.633 |      0.033 |           |
## -----|-----|-----|-----|
##
```



```
##
```

```
CrossTable(x=wine.knn$k5, y=wine.test$quality, prop.chisq = F)
```

```
##
```

```
##
```

```
##      Cell Contents
```

```
## |-----|
## |                N |
## |      N / Row Total |
## |      N / Col Total |
## |      N / Table Total |
## |-----|
```

```
##
```

```
##
```

```
## Total Observations in Table:  1470
```

```
##
```

```
##
```

```
##      | wine.test$quality
## wine.knn$k5 |      Low |      Medium |      High | Row Total |
## -----|-----|-----|-----|-----|
##      Low |      283 |      139 |        1 |      423 |
##      |      0.669 |      0.329 |      0.002 |      0.288 |
##      |      0.576 |      0.149 |      0.020 |      |
##      |      0.193 |      0.095 |      0.001 |      |
## -----|-----|-----|-----|-----|
##      Medium |      206 |      782 |      44 |      1032 |
##      |      0.200 |      0.758 |      0.043 |      0.702 |
##      |      0.420 |      0.841 |      0.898 |      |
##      |      0.140 |      0.532 |      0.030 |      |
## -----|-----|-----|-----|-----|
##      High |        2 |        9 |        4 |        15 |
##      |      0.133 |      0.600 |      0.267 |      0.010 |
##      |      0.004 |      0.010 |      0.082 |      |
##      |      0.001 |      0.006 |      0.003 |      |
## -----|-----|-----|-----|-----|
## Column Total |      491 |      930 |      49 |      1470 |
##      |      0.334 |      0.633 |      0.033 |      |
## -----|-----|-----|-----|-----|
```

```
##
```

```
##
```

```
CrossTable(x=wine.knn$k7, y=wine.test$quality, prop.chisq = F)
```

```
##
```

```
##
```

```
##      Cell Contents
```

```
## |-----|
## |                N |
## |      N / Row Total |
## |      N / Col Total |
## |      N / Table Total |
## |-----|
```

```
##
```

```
##
```

```
## Total Observations in Table: 1470
```

```
##
```

```
##
```

```
## | wine.test$quality
```

```
## wine.knn$k7 | Low | Medium | High | Row Total |
```

```
## -----|-----|-----|-----|-----|
```

```
## Low | 279 | 132 | 0 | 411 |
```

```
## | 0.679 | 0.321 | 0.000 | 0.280 |
```

```
## | 0.568 | 0.142 | 0.000 | |
```

```
## | 0.190 | 0.090 | 0.000 | |
```

```
## -----|-----|-----|-----|-----|
```

```
## Medium | 211 | 791 | 49 | 1051 |
```

```
## | 0.201 | 0.753 | 0.047 | 0.715 |
```

```
## | 0.430 | 0.851 | 1.000 | |
```

```
## | 0.144 | 0.538 | 0.033 | |
```

```
## -----|-----|-----|-----|-----|
```

```
## High | 1 | 7 | 0 | 8 |
```

```
## | 0.125 | 0.875 | 0.000 | 0.005 |
```

```
## | 0.002 | 0.008 | 0.000 | |
```

```
## | 0.001 | 0.005 | 0.000 | |
```

```
## -----|-----|-----|-----|-----|
```

```
## Column Total | 491 | 930 | 49 | 1470 |
```

```
## | 0.334 | 0.633 | 0.033 | |
```

```
## -----|-----|-----|-----|-----|
```

```
##
```

```
##
```

```
# Different values of k produce very similar results with the best being a toss up between k=3 & k=5  
# I think that overall this model performed acceptably. It would be interesting to compare to other  
# models (Naive Bayes, Decesion Tree, etc) and see how it compares in terms of accuracy.
```

```
#
```

```
# I think another important concept that wasn't included in this assignment is the concept of "research  
# and "metrics of success". For example, below I listed the True Positive rates for the three levels: l  
# medium and high. However, perhaps True Positives aren't the most important metrics for your research  
# Maybe False Positives are a bigger issue, or False Negatives. This is context specific but extremely  
# with these types of models. For example, look at the classification of High Quality wines: The highest  
# is approximately 26%. That might not be considered very high depending on your research question.
```

```
#
```

```
#
```

```
#
```

```
# TP rates (k=3): 59.8%, 74.9%, 26.3% for Low, medium & high respectively
```

```
#
```

```
# Note above values are approximate and will change with each run of code unless pseudo-random number g  
# is held constant with set.seed(). I chose not to do this as I was interested to see how different ite  
# would affect the results.
```

## Extra

### Visualization

I tried to plot the KNN output of this assignment visually and ran into some trouble I still believe this to be a great exercise but perhaps not ideal for this particular assignment.

## Some resources:

- <https://cran.r-project.org/web/packages/ElemStatLearn/ElemStatLearn.pdf> # pg 8: mixture.example examples
- <https://stats.stackexchange.com/questions/21572/how-to-plot-decision-boundary-of-a-k-nearest-neighbor-classifier-from-21602#21602>
- <https://stackoverflow.com/questions/31234621/variation-on-how-to-plot-decision-boundary-of-a-k-nearest-neighbor-classifier>

I adapted my code as much as possible from the examples and discussion on Stack Overflow but in the end had to settle for a relatively straightforward cluster plot based on kNN algorithm. I think this could be improved with dimensionality reduction: There are 12 dimensions in this data set so reducing them to two or three principal components would perhaps make for a better visual and also contribute in general to the analysis. I chose to plot pH vs. free.sulfur.dioxide because these are the least correlated factors in the dataset.

```
require(dplyr)
```

```
## Loading required package: dplyr
```

```
##
```

```
## Attaching package: 'dplyr'
```

```
## The following object is masked from 'package:reshape':
```

```
##
```

```
##      rename
```

```
## The following objects are masked from 'package:stats':
```

```
##
```

```
##      filter, lag
```

```
## The following objects are masked from 'package:base':
```

```
##
```

```
##      intersect, setdiff, setequal, union
```

```
classif <- wine.knn$k3
```

```
prob <- attr(classif, "prob")
```

```
wine.gf.df <- bind_rows(mutate(wine.test[, -12],
```

```
    prob=prob,
```

```
    cls="Low",
```

```
    prob_cls=ifelse(classif==cls, 1, 0)),
```

```
  mutate(wine.test[, -12],
```

```
    prob=prob,
```

```
    cls="Medium",
```

```
    prob_cls=ifelse(classif==cls, 1, 0)),
```

```
  mutate(wine.test[, -12],
```

```
    prob=prob,
```

```
    cls="High",
```

```
    prob_cls=ifelse(classif==cls, 1, 0)))
```

```
# wine.gf.df.unq <- wine.gf.df[!duplicated(wine.gf.df[, c('pH', 'free.sulfur.dioxide')]), ]
```

```
require(ggplot2)
```

```
## Loading required package: ggplot2
```


```
ggplot(wine.gf.df) +
```

```
  geom_point(aes(x=pH, y=free.sulfur.dioxide, col=cls),
```

```
    data=mutate(wine.test[, -12], cls=classif),
```

```
    size=1.2) +
```

```
# geom_contour(aes(x=pH, y=free.sulfur.dioxide, z=prob_cls, group=cls, color=cls),
#             bins=2,
#             data=wine.gf.df.unq) +
geom_point(aes(x=x, y=y, col=cls),
           size=3,
           data=data.frame(x=wine.train$pH, y=wine.train$free.sulfur.dioxide, cls=wine.train$quality))
```




## PCA

I decided to spend some more time on this dataset and applied PCA. My goal was the same as above; to organize the data in such a way to be able to create an effective visual. I was unsuccessful in that regard but did notice that the model performs better, with regards to the TP Rate, than without doing PCA.

```
wine.pca <- prcomp(wine.df[, -12], scale. = T, center = T)
screeplot(wine.pca, type="lines")
```

## wine.pca



```
wine.pca.knn <- knn(train=wine.pca$x[wine.train_index,], test=wine.pca$x[-wine.train_index,], cl=wine.d
CrossTable(x=wine.pca.knn, y=wine.df[-wine.train_index,'quality'], prop.chisq = F)
```

```
##
##
##      Cell Contents
## |-----|
## |               N |
## |      N / Row Total |
## |      N / Col Total |
## |      N / Table Total |
## |-----|
##
##
## Total Observations in Table:  1470
##
##
##      | wine.df[-wine.train_index, "quality"]
## wine.pca.knn |      Low |      Medium |      High | Row Total |
## -----|-----|-----|-----|-----|
##      Low |      277 |      149 |      1 |      427 |
##      |      0.649 |      0.349 |      0.002 |      0.290 |
##      |      0.564 |      0.160 |      0.020 |      |
##      |      0.188 |      0.101 |      0.001 |      |
## -----|-----|-----|-----|-----|
```

##	Medium		209		760		38		1007	
##			0.208		0.755		0.038		0.685	
##			0.426		0.817		0.776			
##			0.142		0.517		0.026			
##	-----		-----		-----		-----		-----	
##	High		5		21		10		36	
##			0.139		0.583		0.278		0.024	
##			0.010		0.023		0.204			
##			0.003		0.014		0.007			
##	-----		-----		-----		-----		-----	
##	Column Total		491		930		49		1470	
##			0.334		0.633		0.033			
##	-----		-----		-----		-----		-----	
##										
##										

```


prob <- attr(wine.pca.knn, "prob")
test <- as.data.frame(wine.pca$x)

wine.pca.gf.df <- bind_rows(mutate(test[-wine.train_index,],
                                   prob=prob,
                                   cls="Low",
                                   prob_cls=ifelse(wine.pca.knn==cls, 1, 0)),
                             mutate(test[-wine.train_index,],
                                   prob=prob,
                                   cls="Medium",
                                   prob_cls=ifelse(wine.pca.knn==cls, 1, 0)),
                             mutate(test[-wine.train_index,],
                                   prob=prob,
                                   cls="High",
                                   prob_cls=ifelse(wine.pca.knn==cls, 1, 0)))

wine.pca.gf.df$cols <- vector(length=nrow(wine.pca.gf.df))
wine.pca.gf.df$cols[wine.pca.gf.df$cls == 'Low'] <- 'c'
wine.pca.gf.df$cols[wine.pca.gf.df$cls == 'Medium'] <- 's'
wine.pca.gf.df$cols[wine.pca.gf.df$cls == 'High'] <- 'v'

ggplot(wine.gf.df) +
  geom_point(aes(x=PC1, y=PC2, col=cls),
             data=mutate(test[-wine.train_index,], cls=wine.pca.knn),
             size=1.2) +
  # geom_contour(aes(x=PC1, y=PC2, z=prob_cls, group=cls, color=cols),
  #             bins=2,
  #             data=wine.pca.gf.df) +
  geom_point(aes(x=x, y=y, col=cls),
             size=3,
             data=data.frame(x=wine.pca$x[,1], y=wine.pca$x[,2], cls=wine.df$quality))

```



Save your R codes in an RMD file. Send your RMD and PDF files to the course shell.